

Divide and Conquer

Maximum Contiguous Subarray Problem (MCS)

Divide: divide a given problem into two or more subproblems (ideally of approximately equal size)

Conquer: solving each subproblem (directly if small enough or recursively)

Combine: combining the solution of the subproblems into a global solution

DC Examples:

Maximum Contiguous Subarray Counting Inversions Integer Multiplication

Polynomial Multiplication Quicksort and Partition Deterministic and Randomized Selection

MCS Problem

Problem: Find the span of years in which a company earned the most

Input: An array of reals $A[1 \dots n]$

The value of subarray $A[i \dots j]$ is: $V(i, j) = \sum_{x=i}^j A(x)$

Find $i \leq j$, such that $V(i, j)$ is maximized.

A Brute Force Algorithm:

Calculate the value of $V(i, j)$ for each pair $i \leq j$ and return the maximum value

Time Complexity: $O(n^3)$

A Data-reused Algorithm:

Don't need to calculate each $V(i, j)$ from scratch

Exploit the fact: $V(i, j) = \sum_{x=i}^j A[x] = V(i, j-1) + A[j]$

Time Complexity: $O(n^2)$

A divide-and-conquer Algorithm:

$MCS = \max(S1, S2, A)$

S1: the MCS in $A[1 \dots m]$ S2: the MCS in $A[m+1 \dots n]$ A: the MCS across the cut

$A = A1 \cup A2$

A1: MCS among contiguous subarray ending at $A[m]$

A2: MCS among contiguous subarray starting at $A[m+1]$

$T(n) = 2T(n/2) + n$

Time Complexity: $O(n \log n)$

Counting Inversions Problem

A Brute Force Algorithm:

List each pair $i < j$ and count the inversions

Time Complexity: $O(n^2)$

Divide-and-conquer Algorithm:

Divide: separate list into two halves A and B

Conquer: recursively count inversions in each list

Combine: count inversions (a, b) with $a \in A$ and $b \in B$

Return summation of above three counts

In the course of combination, assuming that A and B are sorted:

Scan A and B from left to right

Compare a_i and b_j ,

If $a_i < b_j$, then a_i is not inverted with any element left in B

If $a_i > b_j$, then b_j is inverted with every element right in A

Function Merge-and-Count(A, B) can be executed in $O(n)$ times where n is the number of elements in A and B

$$T(n) = 2T(n/2) + O(n)$$

The Sort-and-Count algorithm counts the number of inversions in a permutation of size n in $O(n \log n)$ time

The Polynomial Multiplication Problem

What do we need to compute exactly?

$$A(x) = \sum_{i=0}^n a_i x^i$$

$$B(x) = \sum_{i=0}^m b_i x^i$$

$$C(x) = A(x)B(x) = \sum_{k=0}^{n+m} c_k x^k$$

Then

$$c_k = \sum_{0 \leq i \leq n, 0 \leq j \leq m, i+j=k} a_i b_j, \text{ for all } 0 \leq k \leq n+m$$

The vector C is the convolution of the vectors A and B (A major problem in digital signal processing)

Brute Force Algorithm:

Time Complexity: $O(n^2)$

The first Divide-and-Conquer Algorithm:

Divide: the original problem is divided into 4 problems of input size $n/2$

Conquer: compute $A_0(x)B_0(x), A_0(x)B_1(x), A_1(x)B_0(x), A_1(x)B_1(x)$ by recursively calling the algorithm 4 times

Combine: add the four polynomials $O(n)$ times

$$T(n) = 4T(n/2) + n$$

Time Complexity: $O(n^2)$ Same order as the brute force algorithm

An improved divide-and-conquer algorithm:

What we really need are the following 3 terms:

$$A_0(x)B_0(x), (A_0(x)B_1(x) + A_1(x)B_0(x)), A_1(x)B_1(x)$$

The three terms can be obtained by using only three multiplications:

$$A_0(x)B_0(x), A_1(x)B_1(x), (A_0(x) + A_1(x))(B_0(x) + B_1(x))$$

$$T(n) = 3T(n/2) + n$$

Time Complexity: $O(n^{\log 3})$

Analysis of the D-C algorithm

The divide-and-conquer approach does not always give you the best solution

There is actually an $O(n \log n)$ solution to the polynomial multiplication problem

It involves using the Fast Fourier Transform algorithm as a subroutine

The FFT is another classic divide-and-conquer algorithm

The idea of using 3 multiplications instead of 4 is used in large-integer multiplications

Quicksort Problem

Partition:

Given: An array of numbers

Partition: Rearrange the array $A[p \dots r]$ in place into two (possibly empty) subarrays $A[p \dots q-1]$ and $A[q+1 \dots r]$ such that: $A[u] < A[q] < A[v]$ for any $p \leq u \leq q-1$ and $q+1 \leq v \leq r$
 x is called the pivot

Use $A[r]$ as the pivot, and grow partition from left to right

Initially, $(i, j) = (p-1, p)$

Increase j by 1 each time to find a place for $A[j]$

At the same time increase i when necessary

Stops when $j = r$

if $A[j] > x$: $j += 1$

if $A[j] < x$: $i += 1, A[i] \leftrightarrow A[j], j += 1$

Time Complexity: $O(r - p)$

Randomized partition:

In the algorithm Randomized-Partition(A, p, r), we randomly choose an j , $p \leq j \leq r$, use $A[j]$ as pivot

Idea is that if we choose randomly, then the chance that we get unlucky every time is extremely low

Average case analysis:

used for deterministic algorithms

assume the input is chosen randomly from same distribution

Expected case analysis:

used for randomized algorithms

need to work for any given input

Expected Running time of Randomized-Quicksort is $O(n \cdot \log n)$

Priority Queue:

Priority Queue is an abstract data structure that supports two operations:

Insert: inserts the new element into the queue

Extract-Min: removes and returns the smallest element from the queue

Possible implementations:

Unsorted List:

Insert in $O(1)$ time

Extract-Min in $O(n)$ time

Sorted Array:

Insert in $O(n)$ time

Extract-Min in $O(1)$ time

Question: is there any data structure that supports both these operations in $O(\log n)$ time?

Binary Heap:

Heap-order Property: The value of a node is at least the value of its parent

If the heap-order property is maintained, heaps support the priority queue operation in $O(\log n)$ time

Insertion:

Add the new element to the next available position at the lowest level

Restore the min-heap property if violated

General strategy is bubble up: if the parent of the element is larger than the element, then interchange the parent with its child

After each swap, the min-heap property is satisfied for the subtree rooted at the new element

Time Complexity = $O(\text{height}) = O(\log n)$

Extract-Min:

Copy the last element to the root (overwrite the minimum element stored there)

Restore the min-heap property by bubble down: if the element is larger than either of its children, then interchange it with the smaller of its children

Time Complexity = $O(\text{height}) = O(\log n)$

Heapsort

Build a binary heap of n elements

The minimum element is at the top of the heap

Insert n elements one by one: $O(\log n)$

Perform n Extract-Min operations

The elements are extracted in sorted order

Each Extract-Min operation takes $O(\log n)$ time: $O(n \cdot \log n)$

Total Time Complexity: $O(n \cdot \log n)$

Heapsort is as efficient as merge sort and quicksort

Lower Bound for Comparison-based Sorting:

Objective:

All sorting algorithms seen so far are based on comparing elements

Insertion sort has worst-case running time $\Theta(n^2)$, while the others have worst-case running time $\Theta(n \cdot \log n)$

Decision-Tree Model:

A decision tree can model the execution of any comparison-based sorting algorithm

Worst-case running time = height of tree

Lower Bound for Sorting:

Any comparison-based sorting algorithm require $\Omega(n \cdot \log n)$ comparisons

Proof:

A decision tree to sort n elements must have at least $n!$ leaves, since there are $n!$ possible orderings

A binary tree of height h has at most 2^h leaves

Thus, $n! \leq 2^h \implies h \geq \log n! = \Omega(n \log n)$